

A Machine Learning Approach to Analyze Income Determinants in the Mexican Labor Market: The Role of Education (2000–2020)

École Polytechnique, ENSAE, Télécom Paris

Anahi Reyes Miguel¹

April 2025

1. Introduction

The relationship between education and income has long been a central theme in labor economics (Becker, 1964; Mincer, 1974). While many studies have confirmed that educational attainment correlates with higher earnings, this association can be influenced by labor market structures and social inequalities. In the case of Mexico, income distribution is shaped not only by education but also by informality, regional disparities (Hausmann et al., 2021), and occupational mismatches Quinn and Rubb (2006).

Recent advances in machine learning (ML) offer new tools to analyze income dynamics by flexibly capturing non-linearities, interactions, and temporal changes. Unlike traditional econometric approaches that focus on causal inference, ML emphasizes predictive performance and provides tools—like SHAP values—for interpreting feature contributions. These methods are particularly suited to understanding how the role of education in predicting income has evolved over time.

Several recent studies support the use of ML to study income determinants. Herrera et al. (2023) use XGBoost and neural networks to show how education and digital access jointly shape income distribution in Brazil using survey data. They rely on SHAP values to reveal complex feature interactions, highlighting the importance of interpretability in ML-driven policy research.

Matkowski (2021) compares traditional and ML-based income prediction models using U.S. Census data. His work demonstrates that Gradient Boosting and Random Forest outperform OLS, with education consistently emerging as a key predictor—even in the

1. anahi.reyes-miguel@polytechnique.edu

presence of high-dimensional features.

Finally, Gomez-Cravioto et al. (2022) predict alumni income in Mexico using a combination of econometric and ML models using survey data. Education-related features such as GPA and field of study rank among the most important predictors, with SHAP values used to interpret results. Despite its focus on a specific (elite) university population, the study underscores the relevance of ML for understanding wage trajectories.

Building on these insights, this study applies interpretable ML models to harmonized Mexican census data from 2000 to 2020 to assess the evolving role of education in income prediction. The central research question is:

Has the importance of education in predicting income increased over time (2000–2020), relative to other factors such as age, gender, occupation, or local labor market conditions?

Using models such as Linear Regression, Lasso, Ridge, Random Forest, and Gradient Boosting, this project evaluates the predictive contribution of education alongside other sociodemographic and labor-related variables. SHAP values are used to enhance model transparency and trace how the relevance of education has changed over the past two decades.

2. Data

2.1 Dataset Description

This study draws on harmonized census microdata compiled by Banco de México, which includes individual-level information from the extended questionnaires of the 2000, 2010, and 2020 Population Censuses, as well as the 2015 Intercensal Count. It includes sociodemographic and labor-related variables, allowing for consistent comparison over time. It offers complete national coverage and municipal-level representativeness. Due to computational and processing limitations, the analysis was conducted on a simple random sample of 72,000 observations drawn from the full dataset.¹

The target variable for prediction is the **monthly labor income**, denoted as `INGRESO`. The analysis focuses on individual-level characteristics and contextual labor market factors that influence income levels across the Mexican population between 2000 and 2020.

1. See the SIDIE Datasets documentation and methodological notes from Banco de México, available at: <https://www.banxico.org.mx/DataSetsWeb/dataset?ruta=LLM&idioma=en>.

The main explanatory variables reflect commonly studied determinants of earnings in the literature (Hausmann et al., 2021), which are the following:

- **Educational attainment** (ESCOLARIDAD_ACUMULADA): Accumulated years of formal schooling.
- **Age** (EDAD): Continuous variable representing the individual’s age in years.
- **Gender** (LLAVE_SEXO): Binary indicator for male (1) and female (2).
- **Ethnic identification** (LLAVE_PERTEINDIGENA, LLAVE_AFRODES): Categorical variables indicating self-reported Indigenous or Afrodescendant identity.
- **Marital status** (LLAVE_SITUACONYUGAL): Categorical variable distinguishing legal, religious, and informal unions, as well as single, separated, or widowed individuals.
- **Religion** (LLAVE_RELIGION): Self-reported religious affiliation, including both institutional and informal beliefs.
- **Primary activity** (LLAVE_ACTPRIMARIA): Declared main occupation or economic activity.
- **Commute time to work** (LLAVE_TIETRASLADO_TRABAJO): Reported travel time (in minutes) between home and work.
- **Local labor market** (MERCADO_TRABAJO_LOCAL): Geographic identifier representing the regional labor market context.

2.2 Data processing

To prepare the dataset for analysis, I excluded variables with more than 10% missing values, while the remaining data underwent imputation using median and mode strategies for numerical and categorical variables, respectively. To simplify interpretation and reduce noise from rare categories, two grouped variables were constructed: **RELIGION_GROUPED**, indicating whether an individual is religious or not, and **ETHNIC_MINORITY**, identifying self-reported Indigenous or Afrodescendant status. The sample was restricted to individuals aged 18 and above with strictly positive labor income, thereby focusing on the economically active population and excluding those unlikely to participate meaningfully in the labor market, such as students or informal workers with undeclared earnings.

To address the strong right-skewness in income distribution and improve model stability, a log transformation (**log1p**) was applied to the income variable as shown in Figure 4.

While the dataset is nationally representative, it may underreport informal income. Additionally, excluding individuals with zero income and younger populations introduces selection bias, and the education variable may not reflect differences in quality or skill acquisition.

2.3 Data Exploration

After preprocessing, the final analytical sample includes 15,996 individuals. To ensure fair evaluation, the data is split into a training set (90%, 14,396) for model estimation and a testing set (10%, 1,600) for assessing out-of-sample predictive performance.

Table 1 summarizes the main characteristics of the training data. Individuals have an average of 8.8 years of schooling and are, on average, 37 years old. Reported income is highly dispersed, with a median of 3,440 pesos and a maximum close to 100,000,000 pesos. This level of skewness introduces potential risks for model training, such as overfitting to extreme values and reduced performance on typical income ranges.

Table 1: Descriptive Statistics of Key Numerical Variables

	ESCOLARIDAD_ACUMULADA	EDAD	INGRESO
Count	14,396	14,396	14,396
Mean	8.83	37.19	32,929.40
Std	4.53	13.49	1,666,680.71
Min	0.00	18.00	2.00
25%	6.00	26.00	2,143.00
Median	9.00	35.00	3,440.00
75%	12.00	46.00	6,000.00
Max	23.00	98.00	99,999,999.00

Figure 6 shows the distribution of key demographics. Most individuals are aged 20–50, two-thirds are male, and schooling typically ranges from 6 to 12 years. The majority are married and religious, and about one-third identify as part of an ethnic minority, though many did not specify.

Regarding income distribution, Figure 1 shows that income increases with education across all years, though the slope flattens or is even negative at higher education levels. Differences between years suggest temporal shifts in the returns to education.

The correlation matrix (Figure 7) highlights moderate associations between accumulated education and income (0.45), while other variables show weak or no linear correlation, suggesting low risk of multicollinearity.




Figure 1: Average log-income by accumulated years of education across census years (2000–2020).

Missing values are concentrated in a few categorical variables. Figure 5 shows the top categories for each variable after grouping, helping to assess the distribution of remaining missing or unspecified values.

3. Methodology

3.1 Data Preparation and Transformation

To prepare the data for supervised machine learning, I implemented a series of transformations tailored to the structure and quality of the dataset. These steps ensure compatibility with scikit-learn pipelines and enhance model performance and interpretability Pedregosa et al. (2011).

Handling Missing Values: Missing values were addressed using `SimpleImputer`. For numerical variables, including `ESCOLARIDAD_ACUMULADA`, I imputed missing entries using the median of the training set, an approach that is robust to skewed distributions. For categorical variables, missing values were filled with the most frequent category.

Interaction Variable: To capture non-additive relationships between variables—particularly the evolving impact of education over time—I generated interaction terms by applying a polynomial transformation (degree 2) to the interaction between `ESCOLARIDAD_ACUMULADA` and `ANIO`, allowing the model to learn whether the returns to education have increased or decreased across different years.

Encoding Categorical Variables: Categorical variables were encoded differently de-

pending on their cardinality:

- **Low-cardinality variables** (fewer than or equal to 10 categories), such as gender and marital status, were encoded using *One-Hot Encoding*.
- **High-cardinality variables**, such as `MERCADO_TRABAJO_LOCAL` (774 unique values), were encoded using *Target Encoding*, which replaces each category with the mean log income for that group.

Scaling Numerical Features: All numerical variables were standardized using `StandardScaler`, which transforms them to have zero mean and unit variance. This step is particularly important for regularized linear models such as Ridge and Lasso.

3.2 Modeling Approach

ML models offer key advantages for income prediction: they flexibly capture non-linear relationships, handle high-dimensional data, and detect complex interactions between variables. Unlike traditional econometric approaches, which prioritize causal inference and rely on strong parametric assumptions, ML emphasizes predictive accuracy and allows for the exploration of variable importance without imposing rigid model structures. Prior research supports this methodological choice. For instance, (Mareckova and Pohlmeier, 2017) compare traditional statistical techniques with LASSO-based ML variable selection and find that ML outperforms both human and PCA-based selection in predicting labor outcomes. Their findings highlight how ML methods can enhance model performance, especially when dealing with large numbers of potentially relevant predictors.

In this study, a supervised regression framework is used to predict log-transformed income, enabling the assessment of the evolving role of education relative to other demographic and labor market factors. Supervised models are well-suited for this task as they leverage labeled data to learn mappings between features and outcomes. Regression is chosen over classification since income is a continuous variable, and the focus is on modeling variation in levels of income.

To balance predictive performance with interpretability, I implemented five machine learning models: Linear Regression, Lasso and Ridge (regularized linear models suited for feature selection and multicollinearity), and two tree-based ensemble methods—Random Forest and Gradient Boosting—which capture complex nonlinear relationships. All models were evaluated using 5-fold cross-validation. Regularization parameters for Lasso and Ridge were tuned using `LassoCV` and `RidgeCV`, while tree-based models were optimized via grid search on a representative data subset to reduce computation time.

4. Results

4.1 Model Performance Comparison

Table 2 summarizes the performance of the five machine learning models tested, along with a baseline model that predicts the mean income for all observations. Model evaluation was conducted on the test set using three standard metrics: Mean Squared Error (MSE), Mean Absolute Error (MAE), and the coefficient of determination (R^2). MSE captures the average squared difference between predicted and actual values, giving more weight to large errors. MAE represents the average absolute error and is more interpretable and robust to outliers. Finally, R^2 measures the proportion of variance in the dependent variable that is explained by the model, with higher values indicating better explanatory power. These metrics provide a comprehensive view of each model’s predictive accuracy on unseen data.

Table 2: Model Performance on Test Data

Model	MSE	MAE	R^2
Gradient Boosting	0.4804	0.4980	0.4018
Linear Regression	0.4823	0.4984	0.3995
Ridge Regression	0.4823	0.4984	0.3994
Lasso Regression	0.4898	0.5021	0.3901
Random Forest	0.5526	0.5401	0.3119
Baseline (Mean Predictor)	0.8031	0.6726	0.0000

Regarding the relative performance, as shown in Table 2, Gradient Boosting achieved the best overall results with the lowest MSE (0.4804) and MAE (0.4980), and the highest R^2 (0.4018), meaning it explains about 40% of the variance in log-transformed income. Linear and Ridge regression performed similarly well ($R^2 \approx 0.3995$), while Lasso was slightly less accurate. Random Forest showed lower predictive accuracy ($R^2 = 0.3119$), and the baseline model, which simply predicted the mean, performed worst with $R^2 \approx 0$.

In terms of absolute predictive error, a MAE of 0.50 in log-income terms implies that predictions deviate from the true values by a factor of about 1.65 on average, that is, predicted income is typically 65% higher or lower than the actual value. For instance, a predicted income of 5,000 pesos might correspond to a true value between roughly 3,030 and 8,250 pesos. While this reflects the challenge of modeling income in a highly unequal and segmented labor market, it also demonstrates that machine learning models provide improvements over an average-based approach.

4.2 Feature Importance and Interpretation

To understand which variables contribute most to income prediction, we analyzed feature importance using different strategies in the two best-performing machine learning models: Gradient Boosting and Linear Regression. These models were selected based on their predictive performance (as shown in Table 2) and their complementarity—Linear Regression offers interpretability through model coefficients, while Gradient Boosting provides high predictive accuracy and captures complex nonlinear relationships.

To evaluate the role of individual predictors in each model, we rely on SHAP (SHapley Additive exPlanations) values. SHAP values offer a unified framework for interpreting feature contributions across both linear and nonlinear models by quantifying the marginal impact of each input feature on a given prediction.

Figure 2 presents the SHAP beeswarm plot for the **Gradient Boosting Regressor**. The x-axis shows the SHAP value, which reflects the contribution of each feature to the predicted log-income, while the y-axis ranks the features by average absolute SHAP magnitude. The most influential variable is `MERCADO_TRABAJO_LOCAL`, which was target-encoded using average income within each local labor market. High values of this variable (shown in red) are associated with large positive SHAP values, indicating that individuals located in more prosperous labor markets are predicted to earn more.

The second most important feature is the interaction term `ANIO × ESCOLARIDAD_ACUMULADA`. Here, red-colored points (representing high education in recent years) are concentrated on the right, meaning these values increase the predicted income. This suggests that the returns to education have grown over time: education acquired in more recent years contributes more to income than it did in earlier periods.




Figure 2: SHAP Beeswarm Plot for Gradient Boosting Regressor

Figure 3 displays the SHAP beeswarm plot for the **Linear Regression**. Interestingly, the results diverge from those of Gradient Boosting. In this case, **MERCADO_TRABAJO_LOCAL** ranks only tenth in importance, while the interaction term **ANIO \times ESCOLARIDAD_ACUMULADA** and the main effect **ESCOLARIDAD_ACUMULADA** emerge as the top predictors.

The SHAP pattern for the interaction term implies that, contrary to the nonlinear model, education obtained in more recent years may contribute less to predicted income. This is reflected in the clustering of lower SHAP values even for high feature values. Nevertheless, total years of schooling remains a strong positive contributor. These patterns are consistent with the model coefficients reported in Table 3.

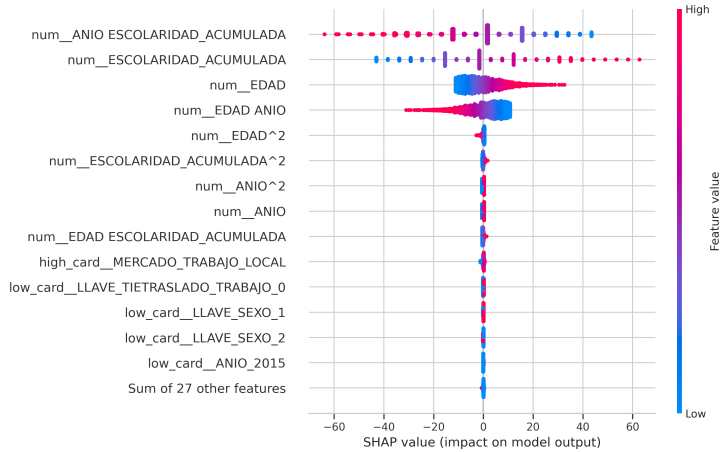


Figure 3: SHAP Beeswarm Plot for Linear Regression

These contrasting patterns reflect the strengths and limitations of each modeling approach. The results also underscore the complexity of the relationship between education and income over time, suggesting that the evolving role of education is not easily captured through a single modeling lens. Rather than offering a definitive answer, this comparison points to the need for further research on how education interacts with time and context to shape economic outcomes. We’re missing granularity of the type of job they do, their degree, and sector — some sector might be better paid than others.

5. Conclusion

This study demonstrates the potential of machine learning approaches to improve income prediction based on individual characteristics. All tested models outperformed a naive baseline that predicted the mean income for every individual, confirming that ML methods can capture meaningful structure in the data beyond what simple averages provide. Among them, Gradient Boosting achieved the highest predictive accuracy, explaining approximately 40% of the variance in log-transformed income.

The contrasting results between Gradient Boosting and Linear Regression highlight the trade-offs between model flexibility and interpretability. Gradient Boosting captured non-linearities and interaction effects, such as the increasing returns to education over time, whereas Linear Regression, with its additive linear structure, offered more transparent coefficients but at the cost of predictive power and oversimplified assumptions. These differences underscore the importance of model selection in interpreting the role of education and other covariates in the formation of income.

However, these insights must be interpreted cautiously given the limitations of the data and computational environment. The analysis is restricted to individuals aged 18 and over with strictly positive reported income, which introduces selection bias and limits the generalizability of findings, particularly in a labor market with widespread informality. Moreover, potential measurement errors or underreporting in income data may reduce model accuracy. Due to computational constraints, models were trained on subsamples, limiting the scope of optimization. With access to more powerful infrastructure, richer modeling strategies and deeper cross-validation could be explored.

This work intentionally focused on individual-level predictors, but future extensions could benefit greatly from incorporating contextual and geographic variables. For example, introducing interaction terms between time and local labor markets (e.g., `ANIO × MERCADO_TRABAJO_LOCAL`) could help uncover evolving regional disparities. Moreover, the inclusion of additional variables such as household composition, housing conditions, or

labor market aggregates could enhance the models' explanatory power. Notably, our analysis is constrained by the lack of detailed information on job types, educational degrees, and employment sectors. Factors that likely play a critical role in shaping income, as certain sectors tend to offer higher wages than others.

While this research does not deliver a definitive model to predict income, it illustrates the relevance of machine learning for exploring labor market dynamics. It emphasizes the importance of combining predictive techniques with thoughtful feature design and interpretability tools, particularly SHAP, to uncover the evolving role of education and geography in income inequality in Mexico. These methods, when carefully applied, can complement traditional economic analysis and inform evidence-based policy.

References

- Becker, G. (1964). *Human Capital*. National Bureau of Economic Research, New York.
- Gomez-Cravioto, D. A., Diaz-Ramos, R. E., Hernandez-Gress, N., and Ceballos, H. G. (2022). Supervised machine learning predictive analytics for alumni income. *Journal of Big Data*, 9(1):11.
- Hausmann, R., Pietrobelli, C., and Santos, M. A. (2021). Place-specific determinants of income gaps: New sub-national evidence from mexico. *World Development*, 146:105566.
- Herrera, G. P., Constantino, M., Su, J.-J., and Naranpanawa, A. (2023). The use of icts and income distribution in brazil: A machine learning explanation using shap values. *Telecommunications Policy*, 47(8):102598.
- Mareckova, J. and Pohlmeier, W. (2017). Noncognitive skills and labor market outcomes: A machine learning approach. In *Beiträge zur Jahrestagung des Vereins für Socialpolitik 2017: Alternative Geld- und Finanzarchitekturen - Session: Treatment Effects*, number G03-V2, Kiel, Hamburg. ZBW - Deutsche Zentralbibliothek für Wirtschaftswissenschaften, Leibniz-Informationszentrum Wirtschaft.
- Matkowski, M. (2021). Prediction of individual income: A machine learning approach. Bachelor's thesis, Bryant University. CC-BY-NC-ND licensed.
- Mincer, J. (1974). *Schooling, Experience, and Wages*. National Bureau of Economic Research, New York.
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau,

- D., Brucher, M., Perrot, M., and Duchesnay, E. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830.
- Quinn, M. A. and Rubb, S. (2006). Mexico’s labor market: The importance of education-occupation matching on wages and productivity in developing countries. *Economics of Education Review*, 25(2):147–156.

A. Additional Material




Figure 4: Distribution of log-transformed income

Table 3: Top 10 Most Impactful Features from Linear Regression.

Feature	Coefficient	Abs. Value
num__ANIO ESCOLARIDAD ACUMULADA	-21.0687	21.0687
num__ESCOLARIDAD ACUMULADA	20.8437	20.8437
num__EDAD	7.3695	7.3695
num__EDAD ANIO	-7.0918	7.0918
high_card__MERCADO TRABAJO LOCAL	1.1408	1.1408
num__EDAD^2	-0.4068	0.4068
num__ESCOLARIDAD ACUMULADA^2	0.3420	0.3420
low_card__LLAVE TIETRASLADO TRABAJO_7	-0.3166	0.3166
num__ANIO	0.3164	0.3164
num__ANIO^2	0.3164	0.3164




Figure 6: Distributions of key demographic characteristics in the training dataset.




Figure 7: Correlation matrix among selected numerical variables.